



Docker Architecture & Docker Networking

Hussaina Begum N.
Staff Engineer, VMware

About me

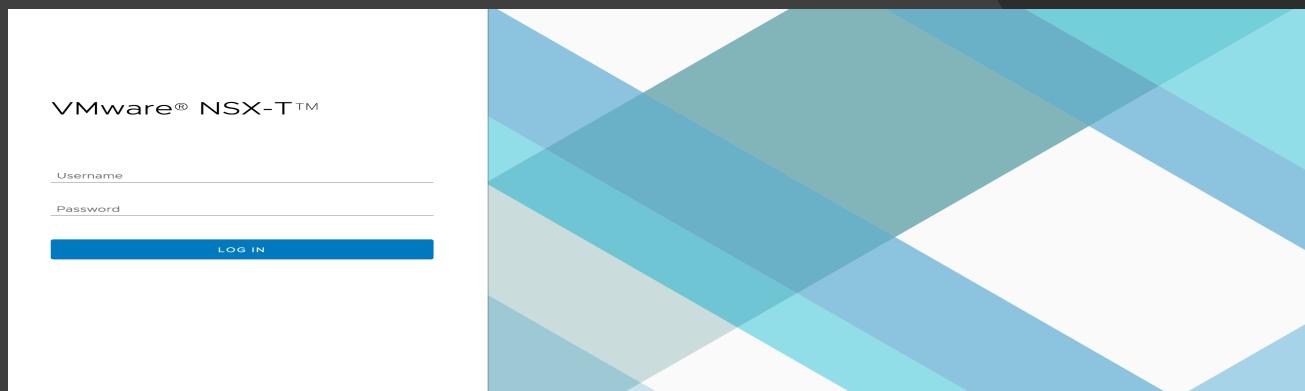


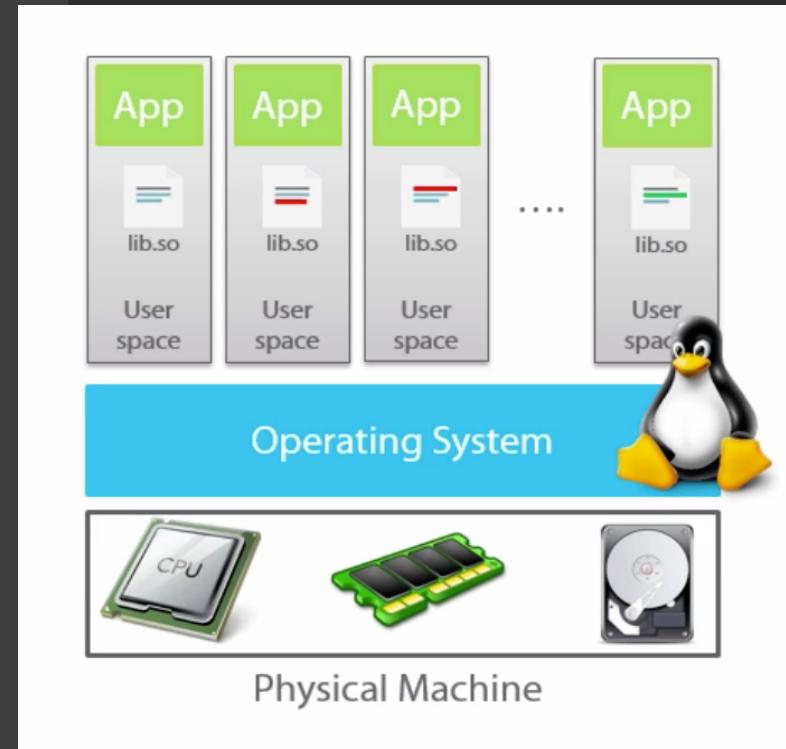
Image copyrights to IBM, Cisco, Citrix & VMware

Agenda

- History of Compute Virtualization
- Linux Containers(LXC)
- Microservices architecture
- What are Dockers ?
- Docker Architecture
- Docker Hands-On
- Docker networking
- Single-Host, Multi-Host networking demos
- Q&A

Long time ago in a galaxy far far away ...

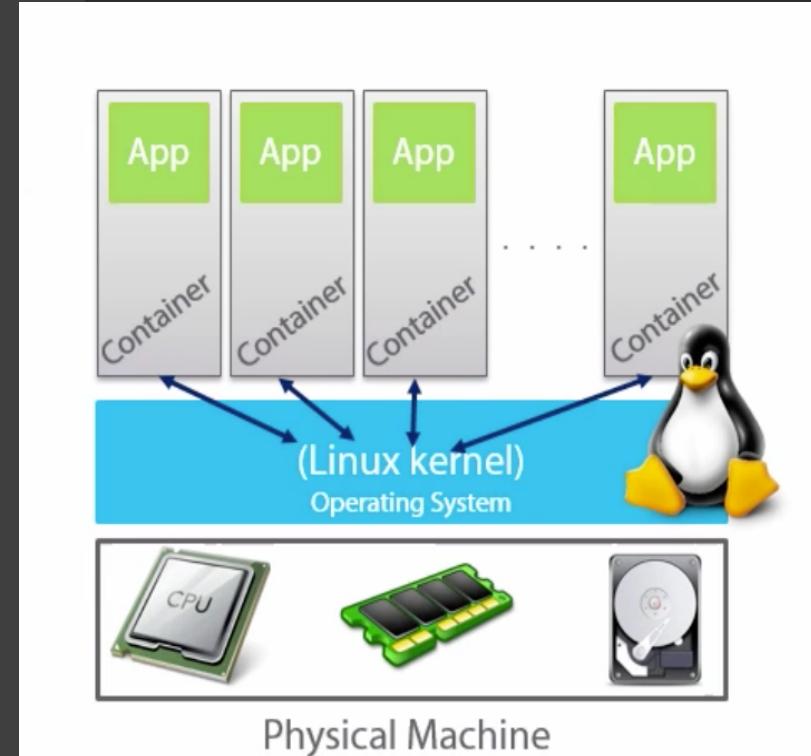
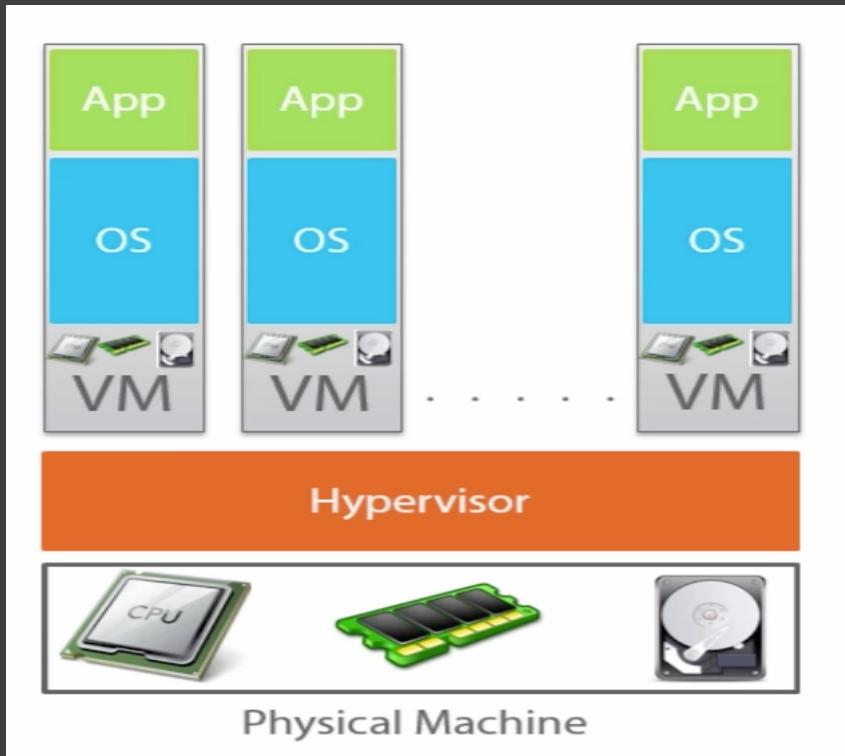
- Physical servers
- Monolithic applications
- Scaling is hard and expensive



@nigelpoulton

History

- VMs vs. containers



@nigelpoulton

Linux Containers (LXC)

- Something in between chroot and VM
- Uses the below constructs to achieve containers
 - Kernel namespaces (ipc, uts, mount, pid, network and user)
 - Apparmor and SELinux profiles
 - Seccomp policies
 - Chroots (using pivot_root)
 - Kernel capabilities
 - Cgroups (control groups)

Linux Containers (LXC) → Docker

- LXC + aufs → dotCloud (dc)
- Changes to LXC broke dc
- dc came up with libcontainer abstraction
- dc -> docker daemon – 2013
- docker daemon is bloated(monolith)
- Open Container Initiative (runtime spec, image spec) - 2015
- docker refactoring

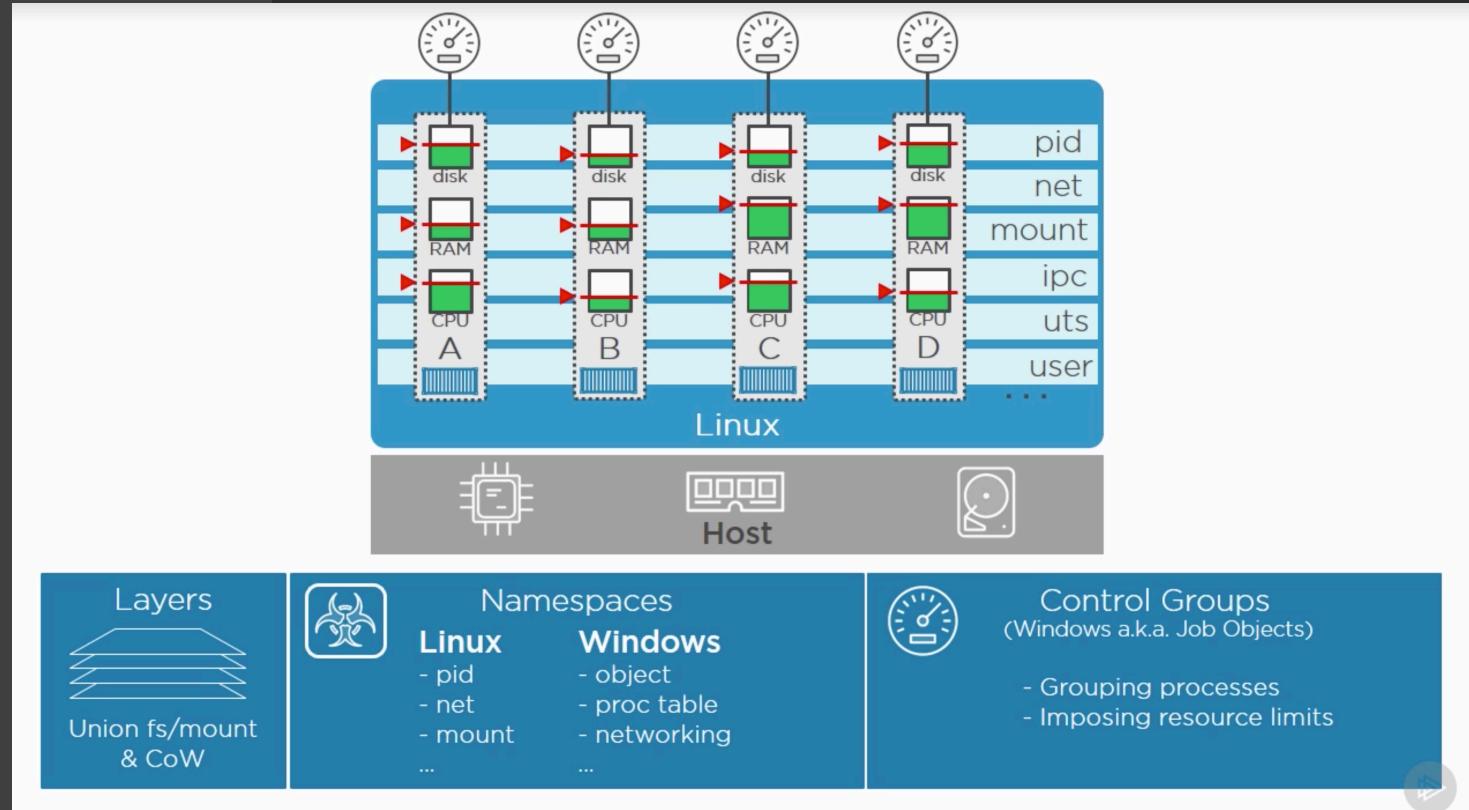
Microservices Architecture

- Monolithic applications are
 - Inflexible
 - Unscalable
 - Unreliable
 - Slow development(same framework for all components)
- Microservices
 - Small autonomous services
 - Independent and loosely coupled components
 - Can scale each component
 - Stateless components
 - Redeployment
 - Small focused, loosely coupled, language independent, bounded context

<https://blog.newrelic.com/technology/microservices-what-they-are-why-to-use-them/>

Dockers

- OS level virtualization
- Shares host kernel
- Multi-layered FS
 - Union FS(overlay2), mount, CoW
- Namespaces
 - pid, net, mnt, ipc, uts, user
- Cgroups
 - Grouping procs, rlimits

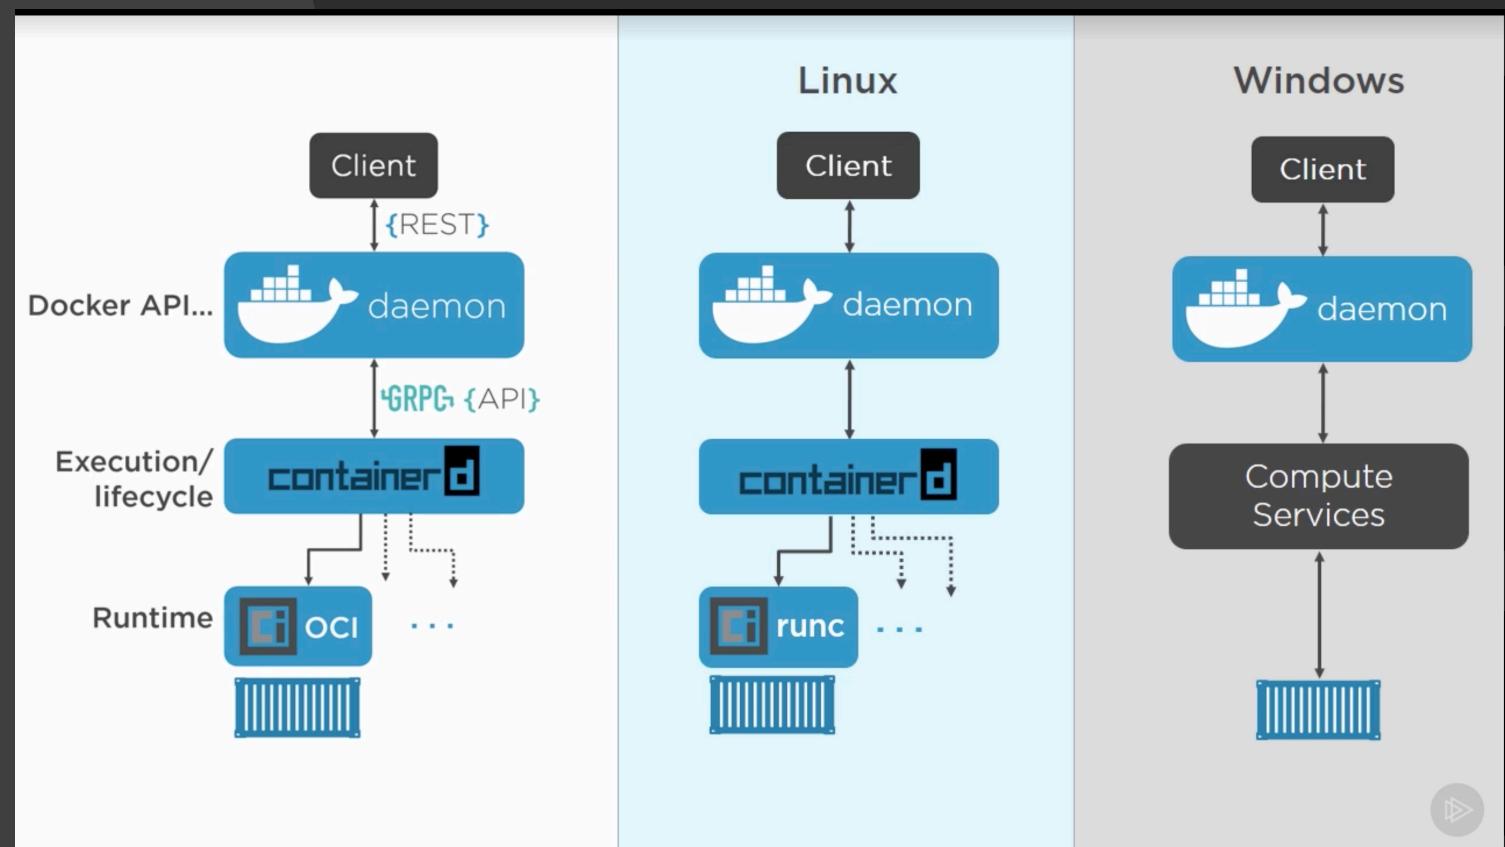


Docker on Windows/Mac

- Docker Toolbox (not a native solution)
 - Virtual Env, boot2docker, Docker Engine, Docker Machine, Docker compose
 - Network, storage is not shared from host
- Mac Native Docker (2016)
 - HW virtualization, Embedded Linux, unikernel, docker engine
- Windows Native Docker(2016)
 - Windows Server based
 - Leverages namespace, cgroup of windows kernel
 - Host kernel is shared
 - Hyper-V based
 - Containers run in highly optimized VM
 - Host kernel is not shared

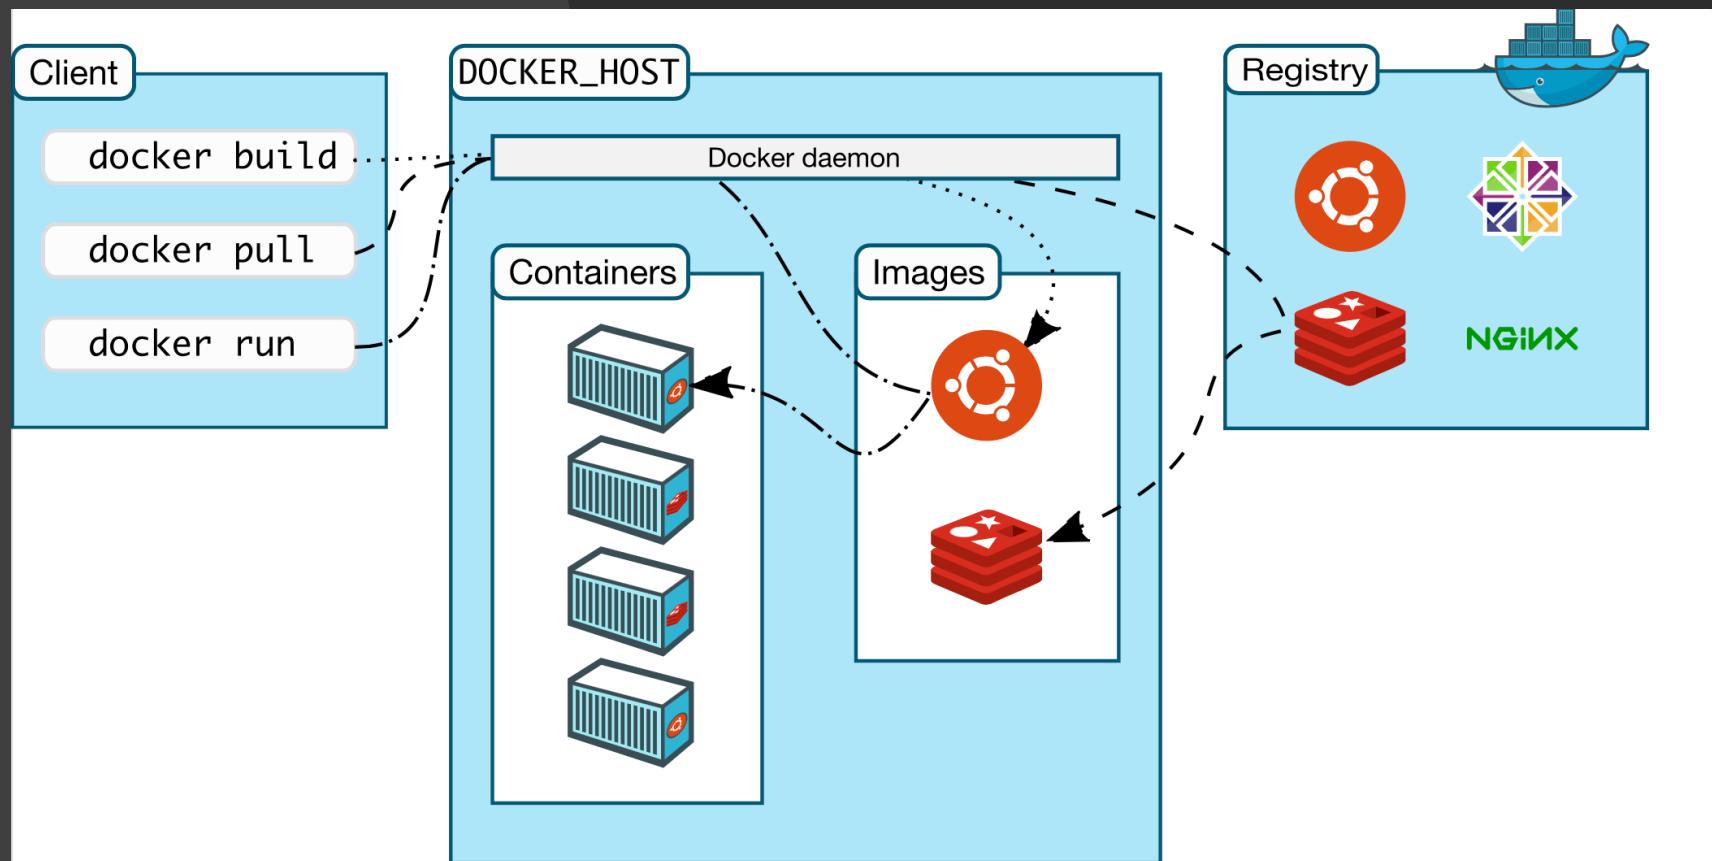
Docker Engine

- API server
- Docker daemon
- containerd
- Runtime
 - runc
 - Rkt – retired
 - Cri-o (for k8s)
 - kata-runtime



Docker workflow

- Image
- Container
- Registry



Dockers – Install/HandsOn

- Docker install on Ubuntu/Mac/Windows
- docker info – lists the docker engine details on the host(client and server versions)
- docker pull <image> – pulls image to docker host
- docker run – starts a new container
- docker stop – stops the container
- docker ps – lists current running containers
- docker images – lists the images on docker host
- docker rmi – removes image from docker host

Dockerfile

- CAPITALIZED instructions
- <INSTRUCTION> - <value>
- FROM
- LABEL
- RUN
- COPY
- WORKDIR
- ENTRYPOINT
- CMD

Container Networking Models

- CNM – by Docker
 - Sandbox
 - Endpoint
 - Network
- CNI – by coreOS (approved by CNCF)
 - CNI Specification - <https://github.com/containernetworking/cni/blob/master/SPEC.md>
 - Plugins
 - Library

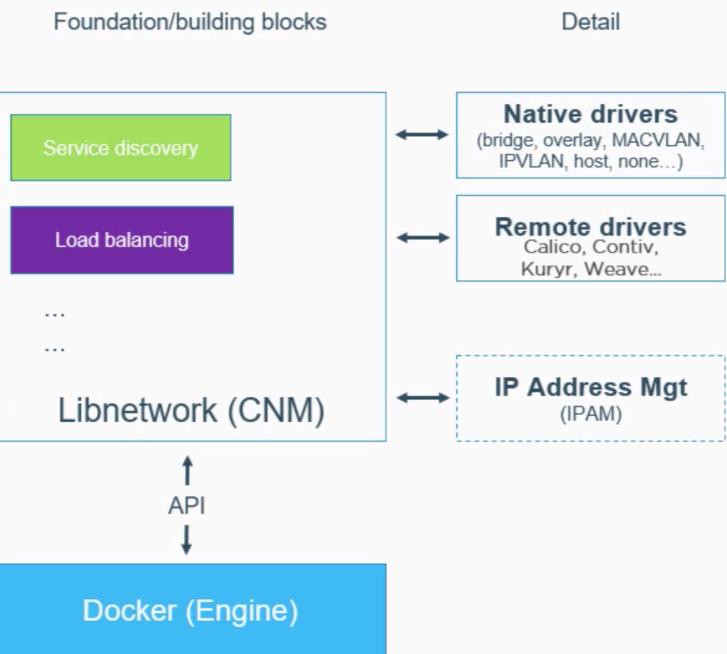
Dockers – Networking

- CNM
 - Sandbox – namespaces
 - Endpoint - interfaces
 - Network – connected endpoints
- Libnetwork
- Drivers
 - host
 - bridge
 - none
 - Remote(overlay, macvlan)
 - plugins

CNM
Master plan/grand design

Libnetwork
De facto implementation of
the CNM

Drivers
Network-specific detail



Dockers – Networking CLIs

- docker network ls
- docker info
- docker network create
- docker network connect
- docker network disconnect

Dockers – Single Host Networking

Drivers that are available on host

- Bridge(default)
- Host
- None

1. Create two containers
2. Create a bridge n/w with subnet
3. Attach bridge n/w to two containers
4. Containers can talk to each other over bridge n/w. i.e, by ip or by container-name

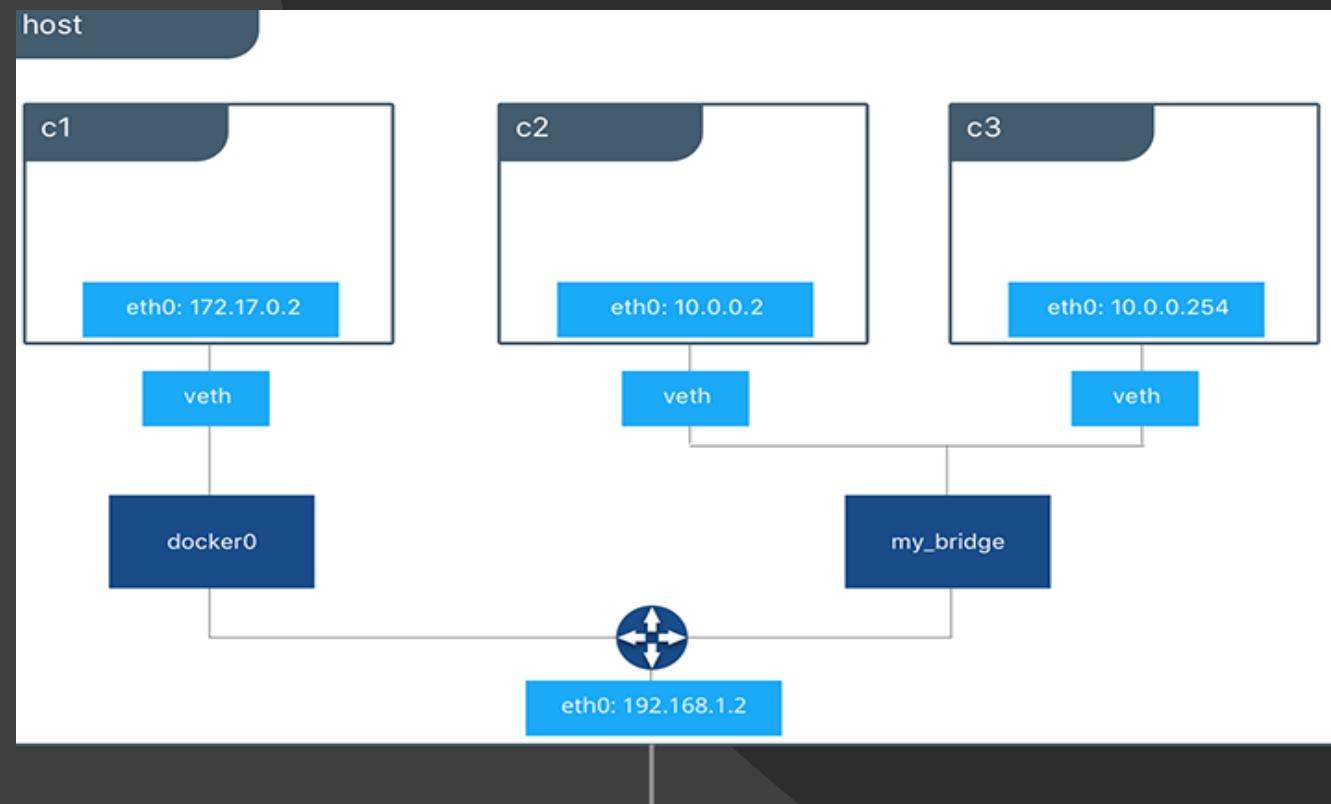
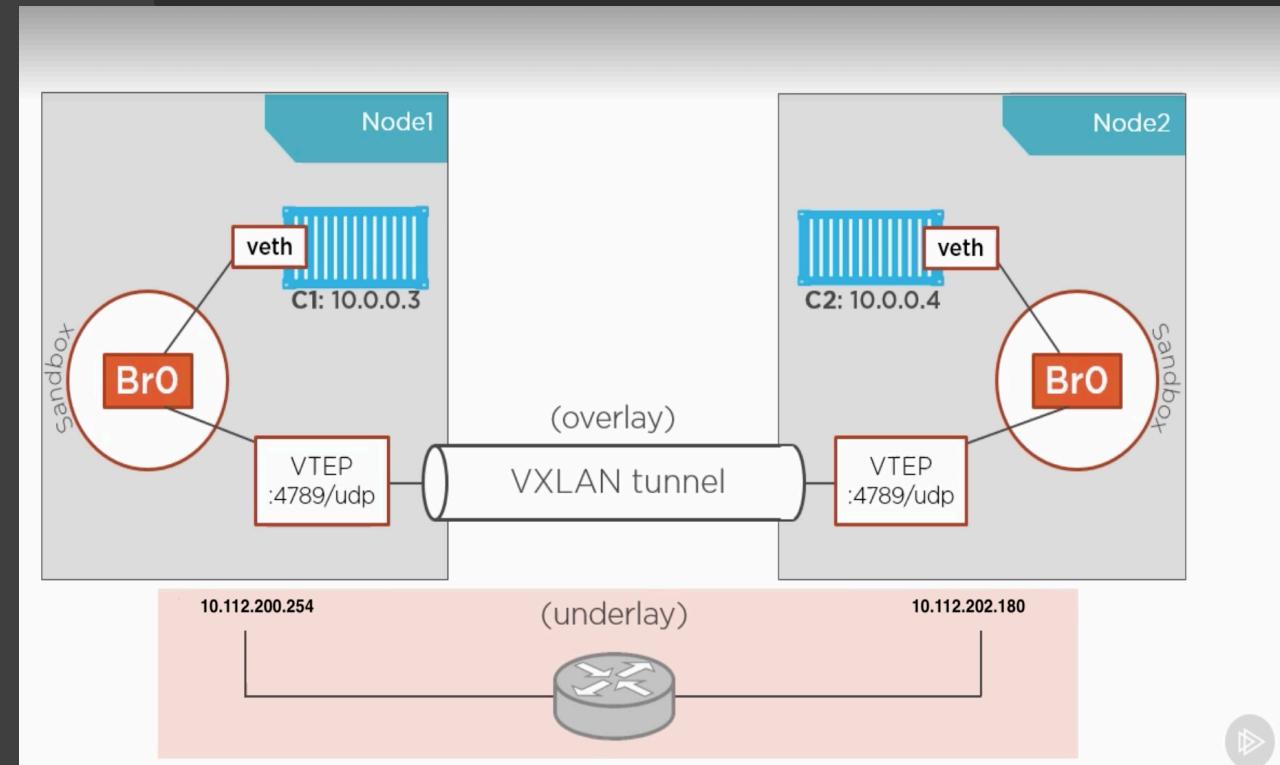


Image @docker.com

Dockers – Multi Host Networking

Overlay driver needs to be used for connectivity.

1. Create two containers on two different nodes
2. Node1 - docker swarm init
3. Node2 - docker swarm join
4. Create network using 'overlay' driver
5. Create a docker service with 2 replicas
6. Containers can talk to each other, using overlay network



Dockers – Orchestration

- docker compose(single host)
 - -link command
- docker swarm (multi-host)
 - swarm init
 - swarm join
- kubernetes

References

- <https://docs.docker.com>
- <https://success.docker.com/article/networking>
- Pluralsight courses
 - <https://app.pluralsight.com/profile/author/nigel-poulton>
 - Docker Deep Dive
 - Docker Networking
- Linux Foundation courses
 - Container Fundamentals (LFS253) -
https://vmware.sabacloud.com/Saba/Web_spf/NA1PRD0121/common/ledetail/cours000000000035007
 - Kubernetes Fundamentals -
https://vmware.sabacloud.com/Saba/Web_spf/NA1PRD0121/common/ledetail/cours000000000035010

Thank you!!
Questions ?